

1.

Part (1)

In logistic_reg, the for loop terminates either iteration times reach max_iterations or when each of element in the gradient goes below the tolerance.

$[w, E_{in}] = \text{logistic_reg}(X_{\text{train}}, y_{\text{train}}, w_{\text{init}}, \text{max_its}, \eta)$

$[E_{\text{class_train}}] = \text{find_test_error}(w, X_{\text{train}}, y_{\text{train}})$

$[E_{\text{test}}] = \text{find_test_error}(w, X_{\text{test}}, y_{\text{test}})$

$\eta = 10^{-5}$, tolerance = 10^{-3} ,

where E_{in} : logistic regression in-sample error

$E_{\text{class_train}}$: classification error for training data set


E_{test} : test error (classification error for testing data set)

max_iterations	10,000	100,000	1,000,000
num_its	10,000	100,000	1,000,000
E_{in}	0.5847	0.4937	0.4354
$E_{\text{class_train}}$	0.3092	0.2237	0.1513
E_{test}	0.3172	0.2069	0.1310
execution_time	4.9 sec	47.9 sec	477.2 sec

What can you say about the generalization properties of the model?

E_{test} would be generally lower than E_{in} , because E_{in} is an estimation bounded by a loose upper bound. Then the actual E_{test} would be lower than the error estimation in training.

Furthermore, if the gradient is not below tolerance, more iterations basically come with lower error rate.



Part (2)

	Best result(max_its = 1 million)	glmfit function
E_{test}	0.1310	0.1103
Execution Time	477.2 sec	0.0403 sec

Part (3)

Scaling the features by subtracting the mean and dividing by the standard deviation for each

of the features: $Z_{train} = \frac{X_{train} - \bar{X}_{train}}{S_{train}}$, $Z_{test} = \frac{X_{test} - \bar{X}_{test}}{S_{test}}$, where S is standard deviation.

Setting tolerance to 10^{-6} . Run iterations until each element of gradient below the tolerance 10^{-6} .

[w, E_{in}] = logistic_reg(Z_train, y_train, w_init, max_its = NONE, η)

[E_{test}] = find_test_error(w, Z_test, y_test)

	$\eta = 10^{-5}$	$\eta = 10^{-4}$	$\eta = 10^{-3}$	$\eta = 10^{-2}$
Number of iterations	23,371,190	2,337,115	233,707	23,367
E_{in}	0.4074	0.4074	0.4074	0.4074
E_{test}	0.1034	0.1034	0.1034	0.1034
Time	12859.8 sec	1086.8 sec	111.9 sec	11.2 sec

	$\eta = 10^{-1}$	$\eta = 1$	$\eta = 3$	$\eta = 5$	$\eta = 7$
Number of iterations	2,332	229	73	41	42
E_{in}	0.4074	0.4074	0.4074	0.4074	0.4074
E_{test}	0.1034	0.1034	0.1034	0.1034	0.1034
Time	1.1 sec	0.11 sec	0.035 sec	0.02 sec	0.022 sec

**when $\eta = 8$ each step is too large, it begins to bouncing around.

2. Problem 3.4

Adaptive Linear Neuron (Adaline) algorithm

(a)

$$e_n(w) = [\max(0, 1 - y_n w^T x_n)]^2 = \begin{cases} 0, & 1 - y_n w^T x_n \leq 0 \\ (y_n w^T x_n)^2, & 1 - y_n w^T x_n > 0 \end{cases}$$

$$\text{when } e_n(w) = 0 \Rightarrow \lim_{(1 - y_n w^T x_n) \rightarrow 0^+} \nabla e_n(w) = 0 = \lim_{(1 - y_n w^T x_n) \rightarrow 0^-} \nabla e_n(w)$$

$$\Rightarrow \nabla e_n(w) = 0$$

$$\text{when } e_n(w) = (y_n w^T x_n)^2$$

$$\lim_{(1-y_n w^T x_n) \rightarrow 0^+} \nabla e_n(w) = \lim_{(1-y_n w^T x_n) \rightarrow 0^+} -2[1 - y_n w^T x_n] y_n x_n = 0$$

$$\lim_{(1-y_n w^T x_n) \rightarrow 0^-} \nabla e_n(w) = \lim_{(1-y_n w^T x_n) \rightarrow 0^-} -2[1 - y_n w^T x_n] y_n x_n = 0$$

$$\Rightarrow \lim_{(1-y_n w^T x_n) \rightarrow 0^+} \nabla e_n(w) = \lim_{(1-y_n w^T x_n) \rightarrow 0^-} \nabla e_n(w) = 0$$

$$\Rightarrow \nabla e_n(w) = 0$$

Therefore, $e_n(w)$ is continuous and differentiable.

(b)

$$\llbracket \text{sign}(w^T x_n) \neq y_n \rrbracket = \begin{cases} 0, & \text{sign}(w^T x_n) = y_n \\ 1, & \text{sign}(w^T x_n) \neq y_n \end{cases}$$

$$e_n(w) = [\max(0, 1 - y_n w^T x_n)]^2 \begin{cases} 0, & 1 - y_n w^T x_n \leq 0 \\ (y_n w^T x_n)^2, & 1 - y_n w^T x_n > 0 \end{cases}$$

$$\text{when } \text{sign}(w^T x_n) = y_n, \quad \llbracket \text{sign}(w^T x_n) \neq y_n \rrbracket = 0$$

$$y_n w^T x_n > 0 \Rightarrow e_n(w) = [0, 1) \geq \llbracket \text{sign}(w^T x_n) \neq y_n \rrbracket$$

$$\text{when } \text{sign}(w^T x_n) \neq y_n, \quad \llbracket \text{sign}(w^T x_n) \neq y_n \rrbracket = 1$$

$$y_n w^T x_n < 0 \Rightarrow e_n(w) > 1 = \llbracket \text{sign}(w^T x_n) \neq y_n \rrbracket$$

Since for each n, $e_n(w) \geq \llbracket \text{sign}(w^T x_n) \neq y_n \rrbracket$, therefore $e_n(w)$ is upper bound of $\llbracket \text{sign}(w^T x_n) \neq y_n \rrbracket$.

$$\Rightarrow \frac{1}{N} \sum_{n=1}^N e_n(w) \geq \frac{1}{N} \sum_{n=1}^N \llbracket \text{sign}(w^T x_n) \neq y_n \rrbracket = E_{in}(w)$$

(c)

$$\begin{aligned} w(t+1) &= w(t) + \eta(1 - y_n w^T x_n) y_n x_n \\ &= w(t) + \eta(y_n - (y_n)^2 w^T x_n) x_n \\ &= w(t) + \eta(y_n - (1) w^T x_n) x_n \\ &= w(t) + \eta(y_n - w^T x_n) x_n \end{aligned}$$

3. Problem 3.19

Pointing out potential problems

(a)

The potential problem is the complexity goes infinite as n goes infinite.

(b)

Same reason as part (a).

The potential problem is the complexity goes infinite as n goes infinite.

(c)

$$\Phi(x) = \{\phi_{i,j}\} \in \mathcal{R}^{(101 \times 101)}$$

Since the transformation's dimension is fixed and finite, there's no potential problem that complexity would go too large.

4. Problem 4.8

$$\begin{aligned} & \begin{cases} w^T w = \sum_{i=1}^n w_i^2 \\ \nabla w = (2w_1, 2w_2, \dots, 2w_n) = 2w \end{cases} \\ E_{aug}(w) &= E_{in}(w) + \lambda w^T \Gamma^T \Gamma w = E_{in}(w) + \lambda w^T w \\ E_{aug}(w) : w(t+1) &\leftarrow w(t) - \eta [\nabla E_{in}(w(t)) + 2\lambda w(t)] \\ &\Rightarrow w(t+1) \leftarrow (1 - 2\eta\lambda)w(t) - \eta \nabla E_{in}(w(t)) \end{aligned}$$

5. Problem 4.25

(a)

$$\text{No, because } E_{out} = E_{in}(g) + O\left(\sqrt{\frac{d}{N}} \ln(N)\right),$$

$$\text{and for each learner, } E_{out} = E_{val}(m) + O\left(\sqrt{\frac{d}{N-K_m}} \ln(N-K_m)\right)$$

Since the second term varies in each m learner, selecting the smallest $E_{val}(m)$ doesn't guarantee the minimum E_{out} .

(b)

$$\text{Yes, because for each learner, } E_{out} = E_{val}(m) + O\left(\sqrt{\frac{d}{N-K_m}} \ln(N-K_m)\right)$$

For every m learner, K_m are all the same, then second term is the same.

Therefore, selecting minimum validation error will lead a minimum E_{out} .

(c)

$$\text{Each } P[E_{out}(m) > E_{val}(m) + \epsilon] \leq e^{-2\epsilon^2 K_m}$$

$$\begin{aligned} P[E_{out} > E_{val} + \epsilon] &\leq \sum_{m=1}^M P[E_{out}(m) > E_{val}(m) + \epsilon] \\ &= \sum_{m=1}^M e^{-2\epsilon^2 K_m} = M e^{-2\epsilon^2 k(\epsilon)} \end{aligned}$$

$$\text{where } k(\epsilon) = \frac{-1}{2\epsilon^2} \ln\left(\frac{1}{M} \sum_{m=1}^M e^{-2\epsilon^2 K_m}\right)$$